



Lecture 3: Convex non-smooth optimisation: the proximal operator, forward-backward splitting

Luca Calatroni

before: CR CNRS, Laboratoire I3S

now: MaLGa, University of Genoa, Italy

Inverse problems in biological imaging

MSc Data Science and Artificial Intelligence

January 21 2025

Motivation for smooth optimisation: Tikhonov regularisation

$$g(x) = \frac{\lambda}{2} \|Bx\|_2^2, \quad B \in \mathbb{R}^{N \times n}$$

- Ridge regularization: $N = n$ and $B = \text{Id}$:

$$x^* \in \arg \min_{x \in \mathbb{R}^n} \frac{1}{2} \|Ax - y\|^2 + \frac{\lambda}{2} \|x\|_2^2$$

Reduces high-values of the image x .

Motivation for smooth optimisation: Tikhonov regularisation

$$g(x) = \frac{\lambda}{2} \|Bx\|_2^2, \quad B \in \mathbb{R}^{N \times n}$$

- **Ridge regularization:** $N = n$ and $B = \text{Id}$:

$$x^* \in \arg \min_{x \in \mathbb{R}^n} \frac{1}{2} \|Ax - y\|^2 + \frac{\lambda}{2} \|x\|_2^2$$

Reduces high-values of the image x .

- **Sobolev regularization** (Tikhonov, Arsenin, '83): $N = 2n$ and $B = D = \begin{pmatrix} D_h \\ D_v \end{pmatrix}$:

$$x^* \in \arg \min_{x \in \mathbb{R}^n} \frac{1}{2} \|Ax - y\|^2 + \frac{\lambda}{2} \|Dx\|_{2,2}^2$$

Reduces high-values of the finite-difference image **gradient** Dx (hence oscillations, but also edge sharpness).



Motivation for smooth optimisation: Tikhonov regularisation

$$g(x) = \frac{\lambda}{2} \|Bx\|_2^2, \quad B \in \mathbb{R}^{N \times n}$$

- **Ridge regularization:** $N = n$ and $B = \text{Id}$:


$$x^* \in \arg \min_{x \in \mathbb{R}^n} \frac{1}{2} \|Ax - y\|^2 + \frac{\lambda}{2} \|x\|_2^2$$

Reduces high-values of the image x .

- **Sobolev regularization** (Tikhonov, Arsenin, '83): $N = 2n$ and $B = D = \begin{pmatrix} D_h \\ D_v \end{pmatrix}$:

$$x^* \in \arg \min_{x \in \mathbb{R}^n} \frac{1}{2} \|Ax - y\|^2 + \frac{\lambda}{2} \|Dx\|_{2,2}^2$$

Reduces high-values of the finite-difference image **gradient** Dx (hence oscillations, but also edge sharpness).



Convexity and continuity

You asked if for a given proper function: $f : \mathbb{R}^n \rightarrow \mathbb{R} \cup \{+\infty\}$ the following holds:

$$f \text{ convex} \quad \Rightarrow \quad \text{continuous}$$

or, in other words, if there is any example of a convex function f which is only l.s.c. but not continuous. Let us consider $f : \mathbb{R}^n \rightarrow \mathbb{R} \cup \{+\infty\}$ as:

$$f(x) = \begin{cases} x & x \geq 0 \\ +\infty & x < 0 \end{cases}.$$

The function is proper, l.s.c. but not continuous. The function is convex since for $\alpha \in [0, 1]$:

- if $x, y \geq 0$, then by definition $f(\alpha x + (1 - \alpha)y) = \alpha x + (1 - \alpha)y = \alpha f(x) + (1 - \alpha)f(y)$. Hence the convexity condition holds with equal sign.
- if $x, y < 0$, $f(x) = f(y) = +\infty$ and $\lambda x + (1 - \alpha)y < 0$ so $f(\lambda x + (1 - \alpha)y) = +\infty$. So you get $+\infty \leq +\infty$.
- let $x \geq 0$ and $y < 0$ (the reverse case is the same). Then $f(x) = x$ and $f(y) = +\infty$. So $\alpha f(x) + (1 - \alpha)f(y) = +\infty$. Now, depending on the size of α :
 - * if $\alpha x + (1 - \alpha)y \geq 0$ then $f(\alpha x + (1 - \alpha)y)$ is finite ($< +\infty$)
 - * if $\alpha x + (1 - \alpha)y < 0$ then $f(\alpha x + (1 - \alpha)y) = +\infty$

Either case there holds $f(\alpha x + (1 - \alpha)y) \leq \alpha f(x) + (1 - \alpha)f(y) = +\infty$.

You can build other similar examples by defining a function to be continuous and convex on a subspace of $\mathbb{R}^n$ and $+\infty$ elsewhere.

Table of contents

1. Subdifferentiability
2. The proximal operator
3. Forward-backward splitting a.k.a. proximal gradient descent
4. Convergence properties of proximal gradient descent
5. Accelerated forward-backward splitting (FISTA)

Subdifferentiability


A preliminary observation

We saw that for f differentiable:

$$f \text{ is convex} \quad \Leftrightarrow \quad (\forall x, y \in \mathbb{R}^n) \quad f(y) \geq \underbrace{f(x) + \nabla f(x)^T (y - x)}_{=:\phi(y;x)}$$

Or, in other words:

- the function $\phi(y; x)$ is an affine lower bound/estimator of f
- the tangent to f is below f at all points.




A preliminary observation

We saw that for f differentiable:

$$f \text{ is convex} \quad \Leftrightarrow \quad (\forall x, y \in \mathbb{R}^n) \quad f(y) \geq \underbrace{f(x) + \nabla f(x)^T (y - x)}_{=:\phi(y;x)}$$

Or, in other words:

- the function $\phi(y; x)$ is an affine lower bound/estimator of f
- the tangent to f is below f at all points.



... what if f is not differentiable (but convex)?

Subdifferential and subgradients

Look at the *non-nice* component (typically, the regularisation) g of the original problem we want to solve:

$$\min_{x \in \mathbb{R}^n} \{F(x) := \cancel{f(x)} + g(x)\},$$

Subdifferential and subgradients

Look at the *non-nice* component (typically, the regularisation) g of the original problem we want to solve:

$$\min_{x \in \mathbb{R}^n} \{F(x) := \cancel{f(x)} + g(x)\},$$

Definition (subdifferentials and subgradients)

Let $g : \mathbb{R}^n \rightarrow \mathbb{R} \cup \{+\infty\}$ be a proper and **convex** function. Then, a vector $p \in \mathbb{R}^n$ is a *subgradient* of g at point $x \in \text{dom}(g)$ iff:

$$g(y) \geq g(x) + \langle p, y - x \rangle = g(x) + p^T(y - x), \quad \forall y \in \text{dom}(g)$$

The set of all subgradients at a point $x \in \mathbb{R}^n$ is called the *subdifferential* of g at x , and it is denoted by:

$$\partial g(x) = \{p \in \mathbb{R}^n : p \text{ is a subgradient of } g \text{ at point } x \in \mathbb{R}^n\}$$

Subdifferential and subgradients

Look at the *non-nice* component (typically, the regularisation) g of the original problem we want to solve:

$$\min_{x \in \mathbb{R}^n} \{F(x) := \cancel{f(x)} + g(x)\},$$

Definition (subdifferentials and subgradients)

Let $g : \mathbb{R}^n \rightarrow \mathbb{R} \cup \{+\infty\}$ be a proper and **convex** function. Then, a vector $p \in \mathbb{R}^n$ is a *subgradient* of g at point $x \in \text{dom}(g)$ iff:

$$g(y) \geq g(x) + \langle p, y - x \rangle = g(x) + p^T(y - x), \quad \forall y \in \text{dom}(g)$$


The set of all subgradients at a point $x \in \mathbb{R}^n$ is called the *subdifferential* of g at x , and it is denoted by:

$$\partial g(x) = \{p \in \mathbb{R}^n : p \text{ is a subgradient of } g \text{ at point } x \in \mathbb{R}^n\}$$

Interpretation:

- $p \in \partial g(x)$ if and only if $\phi(y; x) = g(x) + p^T(y - x)$ is a lower affine bound for g .
- $\partial g(x)$ collects all the **slopes** of tangents passing through x .

In general, $\partial g(x)$ contains many elements (“many derivatives at each point”).



Multiple subgradients at a non-differentiable point x_0 .

However, one can show that if g is differentiable in $x \in \text{dom}(g)$, then:

$$\partial g(x) = \{\nabla g(x)\},$$

i.e. the only element in $\partial g(x)$ is the (classical) gradient of g in x .

Exercise: compute $\partial g(x)$ at all $x \in \mathbb{R}$ for the 1D function $g(x) = |x|$ and provide a graphical representation of the result.

Separable functions

Very often, the n -dimensional function you deal with, can be nicely expressed as the sum of 1D components. For instance, think of:

- **norms** $\|x\|_p^p$, $p \geq 1$: $\|x\|_p^p = \sum_{i=1}^n |x_i|^p$, hence least-square terms
 $\|Ax - y\|_2^2 = \sum_{i=1}^m ((Ax)_i - y_i)^2 \dots$
- **sum of norms**, e.g. $g(x) = \|x\|_1 + \frac{\lambda}{2} \|x\|_2^2 = \sum_{i=1}^n \left(|x_i| + \frac{\lambda}{2} |x_i|^2 \right)$.
- ...

Separable functions

Very often, the n -dimensional function you deal with, can be nicely expressed as the sum of 1D components. For instance, think of:

- **norms** $\|x\|_p^p$, $p \geq 1$: $\|x\|_p^p = \sum_{i=1}^n |x_i|^p$, hence least-square terms $\|Ax - y\|_2^2 = \sum_{i=1}^m ((Ax)_i - y_i)^2 \dots$
- **sum of norms**, e.g. $g(x) = \|x\|_1 + \frac{\lambda}{2} \|x\|_2^2 = \sum_{i=1}^n \left(|x_i| + \frac{\lambda}{2} |x_i|^2 \right)$.
- ...

Definition (separable function)

Let $g : \mathbb{R}^n \rightarrow \mathbb{R} \cup \{+\infty\}$ be a proper, convex and l.s.c. function. We say that g is *separable* if there exist proper, **univariate** convex functions $g_i : \mathbb{R} \rightarrow \mathbb{R} \cup \{+\infty\}$ such that

$$g(x) = \sum_{i=1}^n g_i(x_i), \quad \forall x = (x_1, \dots, x_n) \in \mathbb{R}^n.$$

Proposition (subdifferential of separable functions)

Let $g : \mathbb{R}^n \rightarrow \mathbb{R} \cup \{+\infty\}$ be proper, convex, l.s.c. and separable. Then, for all $x \in \text{dom}(g)$:

$$\partial g(x) = (\partial g_1(x_1)) \times \dots \times (\partial g_n(x_n)).$$

Exercise: compute $\partial g(x)$ at all $x \in \mathbb{R}^n$ of $g : \mathbb{R}^n \rightarrow \mathbb{R}$, $g(x) = \|x\|_1$ using the results above. Then, compute $\partial F(x)$ of $F(x) := \frac{1}{2} \|Ax - y\|_2^2 + \lambda \|x\|_1$.

We have now the tools to introduce **optimality conditions** in the convex, proper, l.s.c. but **non-differentiable** case.

Proposition (optimality conditions for non-smooth, cvx functions)

Let $g : \mathbb{R}^n \rightarrow \mathbb{R} \cup \{+\infty\}$ be proper, convex and l.s.c. ($g \in \Gamma_0(\mathbb{R}^n)$). Then:

$$x^* \in \arg \min_{x \in \mathbb{R}^n} g(x) \quad \Longleftrightarrow \quad 0_{\mathbb{R}^n} \in \partial g(x^*)$$

Optimality conditions

We have now the tools to introduce **optimality conditions** in the convex, proper, l.s.c. but **non-differentiable** case.

Proposition (optimality conditions for non-smooth, cvx functions)

Let $g : \mathbb{R}^n \rightarrow \mathbb{R} \cup \{+\infty\}$ be proper, convex and l.s.c. ($g \in \Gamma_0(\mathbb{R}^n)$). Then:

$$x^* \in \arg \min_{x \in \mathbb{R}^n} g(x) \quad \Longleftrightarrow \quad 0_{\mathbb{R}^n} \in \partial g(x^*)$$

Interpretation:

- The set $\partial g(x^*)$ is, in general, multivalued but as soon as the null vector $0_{\mathbb{R}^n} \in \mathbb{R}^n$ belongs to it, then x^* is a minimiser, and viceversa.
- If g is differentiable, the result reads $0 = \nabla g(x^*)$, which is the Fermat's theorem. The result above is a generalisation to the non-smooth case.

Going towards the solution of the composite problem

Go back to the original, composite, **non-smooth** (since g is) problem:

$$\arg \min_{x \in \mathbb{R}^n} \{F(x) := f(x) + g(x)\}$$

Using the rules above we have:

$$x^* \in \arg \min_{x \in \mathbb{R}^n} F(x) \Leftrightarrow 0 \in \partial F(x^*) = \underbrace{\partial f(x^*)}_{f \text{ is smooth}} + \partial g(x^*) = \{\nabla f(x^*)\} + \partial g(x^*)$$

Going towards the solution of the composite problem

Go back to the original, composite, **non-smooth** (since g is) problem:

$$\arg \min_{x \in \mathbb{R}^n} \{F(x) := f(x) + g(x)\}$$

Using the rules above we have:

$$x^* \in \arg \min_{x \in \mathbb{R}^n} F(x) \Leftrightarrow 0 \in \partial F(x^*) = \underbrace{\partial f(x^*)}_{f \text{ is smooth}} + \partial g(x^*) = \{\nabla f(x^*)\} + \partial g(x^*)$$

Definition (stationary point)

A point $x^* \in \mathbb{R}^n$ verifying:

$$0 \in \{\nabla f(x^*)\} + \partial g(x^*) \quad \Leftrightarrow \quad -\nabla f(x^*) \in \partial g(x^*)$$

is said to be a **stationary point** of the composite functional $F := f + g$.

Example: stationary points in constrained programming problems

Let $C \subset \mathbb{R}^n$ be a closed and convex set. Let us define the **indicator function** of C as:

$$\iota_C(x) := \begin{cases} 0 & \text{if } x \in C \\ +\infty & \text{if } x \notin C \end{cases}$$

The function $\iota_C(x)$ is proper, convex and l.s.c.

Example: stationary points in constrained programming problems

Let $C \subset \mathbb{R}^n$ be a closed and convex set. Let us define the **indicator function** of C as:

$$\iota_C(x) := \begin{cases} 0 & \text{if } x \in C \\ +\infty & \text{if } x \notin C \end{cases}$$

The function $\iota_C(x)$ is proper, convex and l.s.c.

Consider:

$$\arg \min_{x \in C} f(x) \quad (\text{constrained}) \Leftrightarrow \arg \min_{x \in \mathbb{R}^n} f(x) + \iota_C(x) \quad (\text{unconstrained})$$

Example: stationary points in constrained programming problems

Let $C \subset \mathbb{R}^n$ be a closed and convex set. Let us define the **indicator function** of C as:

$$\iota_C(x) := \begin{cases} 0 & \text{if } x \in C \\ +\infty & \text{if } x \notin C \end{cases}$$

The function $\iota_C(x)$ is proper, convex and l.s.c.

Consider:

$$\arg \min_{x \in C} f(x) \quad (\text{constrained}) \Leftrightarrow \arg \min_{x \in \mathbb{R}^n} f(x) + \iota_C(x) \quad (\text{unconstrained})$$

Stationary points $x^* \in C$ need to satisfy:

$$-\nabla f(x^*) \in \partial \iota_C(x^*)$$

By definition of subdifferential we have that $y \in \partial \iota_C(x^*)$ if and only if:

$$\underbrace{\iota_C(z)}_{=0} \geq \underbrace{\iota_C(x^*)}_{=0} + y^T(z - x^*) \quad \text{for all } z \in C$$

Equivalently:

$$y^T(z - x^*) \leq 0 \quad \text{for all } z \in C$$

The set: $N_C(x^*) := \{y \in \mathbb{R}^n : y^T(z - x^*) \leq 0\}$ is the **normal cone** of C at x^* .

The proximal operator

Proximal operator: definition

Crucial tool for the development of non-smooth optimisation algorithms. Relations with activation functions in the context of deep networks ([Combettes, Pesquet, '21](#)).

Proximal operator: definition

Crucial tool for the development of non-smooth optimisation algorithms. Relations with activation functions in the context of deep networks ([Combettes, Pesquet, '21](#)).

Definition (proximal operator)


Let $g : \mathbb{R}^n \rightarrow \mathbb{R} \cup \{+\infty\}$ be a proper, convex, l.s.c. function. Then, the *proximal operator* of g with parameter $\gamma > 0$ is the function $\text{prox}_{\gamma g} : \mathbb{R}^n \rightarrow \mathbb{R}^n$ defined for all $x \in \mathbb{R}^n$ by:

$$\text{prox}_{\gamma g}(x) := \arg \min_{y \in \mathbb{R}^n} \gamma g(y) + \frac{1}{2} \|y - x\|^2 = \arg \min_{y \in \mathbb{R}^n} \underbrace{g(y) + \frac{1}{2\gamma} \|y - x\|^2}_{=: h(y; x)}$$

Remarks:

- For a fixed $x \in \mathbb{R}^n$, the function $y \mapsto h(y; x)$ is the sum of a convex + strictly convex function, hence it is strictly convex. Hence, it has a **unique** minimiser, the **proximal point** $\text{prox}_{\gamma g}(x)$.
- If g is not assumed to be convex, then there may be multiple minimisers. . . [Exercise](#) (if time allows).

Graphical interpretation



Thin black lines: level lines of g . **Thick** black lines: graph of the function. **Blue points**: evaluation points are moved to the **red points** in the minimisation with an amount depending on γ . Note: points are moved to the minimum of the function.

For $\gamma > 0$ and $x \in \mathbb{R}^n$, let $z := \text{prox}_{\gamma g}(x)$. We have:

$$z := \text{prox}_{\gamma g}(x) \quad \Leftrightarrow \quad z = \arg \min_{y \in \mathbb{R}^n} g(y) + \frac{1}{2\gamma} \|y - x\|^2$$

$$\text{(optimality)} \quad \Leftrightarrow \quad 0 \in \partial g(z) + \frac{1}{\gamma}(z - x)$$

$$\text{(rearranging)} \quad \Leftrightarrow \quad x \in z + \gamma \partial g(z)$$

$$\text{(using operators)} \quad \Leftrightarrow \quad x \in (Id + \gamma \partial g)(z) \quad (Id(z) = z)$$

$$\text{(uniqueness)} \quad \Leftrightarrow \quad z = (Id + \gamma \partial g)^{-1}(x)$$

Characterisations of the proximal operator

Lemma (characterisations of prox operator)

Let $g : \mathbb{R}^n \rightarrow \mathbb{R} \cup \{+\infty\}$ be a proper and convex function and $x, z \in \mathbb{R}^n$. The following claims are equivalent ($\gamma = 1$):

- $z = \text{prox}_g(x)$
- $x - z \in \partial g(z)$
- $(x - z)^T(y - z) \leq g(y) - g(z)$ for all $y \in \mathbb{R}^n$

Proof:

By definition:

$$z = \arg \min_{y \in \mathbb{R}^n} g(y) + \frac{1}{2} \|y - x\|^2.$$

By optimality theorem and the sum rule of subdifferential calculus, we get:

$$0 \in \partial g(z) + z - x \Leftrightarrow x - z \in \partial g(z).$$

By applying the definition of subdifferential to the vector $w = x - z$, we get:

$$g(y) \geq g(z) + w^T(y - z) = g(z) + (x - z)^T(y - z), \quad \text{for all } y \in \mathbb{R}^n.$$

Proximal operator and implicit gradient descent

Subgradient descent? For $x_0 \in \mathbb{R}^n$, τ_k suitably chosen

$$x_{k+1} = x_k - \tau_k p_k, \quad \text{where } p_k \in \partial g(x_k)$$

It turns out that this iteration scheme is **very slow** ($O(1/\sqrt{k})$) so not practically used...

Proximal operator and implicit gradient descent

Subgradient descent? For $x_0 \in \mathbb{R}^n$, τ_k suitably chosen

$$x_{k+1} = x_k - \tau_k p_k, \quad \text{where } p_k \in \partial g(x_k)$$

It turns out that this iteration scheme is **very slow** ($O(1/\sqrt{k})$) so not practically used...

A way to improve the speed of convergence is to move from an **explicit** (iteration k) to an **implicit** (iteration $k+1$) update, i.e. considering for $k \geq 0$

$$x_{k+1} = x_k - \tau_k \underbrace{p_k}_{p_{k+1}}, \quad \text{where } p_{k+1} \in \partial g(x_{k+1})$$

Equivalently, we get the **proximal algorithm**

$$x_{k+1} \in x_k - \tau_k \partial g(x_{k+1}) \Leftrightarrow x_k \in x_{k+1} + \tau_k \partial g(x_{k+1})$$

$$x_k \in (Id + \tau_k \partial g)(x_{k+1}) \Leftrightarrow x_{k+1} \ni (Id + \tau_k \partial g)^{-1}(x_k)$$

$$\Leftrightarrow x_{k+1} = \text{prox}_{\tau_k g}(x_k) \quad (\text{by above})$$

Note: same convergence speed as gradient descent, $O(1/k)$.

Computation of proximal operators: examples

Let $C \subset \mathbb{R}^n$ be a closed and convex set. Recall **indicator function** of C as:

$$\iota_C(x) := \begin{cases} 0 & \text{if } x \in C \\ +\infty & \text{if } x \notin C \end{cases}$$

The function $\iota_C(x)$ is proper, convex and l.s.c. Proximal operator?

Computation of proximal operators: examples

Let $C \subset \mathbb{R}^n$ be a closed and convex set. Recall **indicator function** of C as:

$$\iota_C(x) := \begin{cases} 0 & \text{if } x \in C \\ +\infty & \text{if } x \notin C \end{cases}$$

The function $\iota_C(x)$ is proper, convex and l.s.c. Proximal operator?

$$\text{prox}_{\gamma \iota_C}(x) = \arg \min_{y \in \mathbb{R}^n} \iota_C(y) + \frac{1}{2\gamma} \|y - x\|^2 = \arg \min_{y \in C} \frac{1}{2\gamma} \|y - x\|^2 = P_C(x),$$

the (Euclidean) **projection** of x onto C (the closest point $y \in C$ to x).

The notion of prox for functions g more general than ι_C is the reason why the prox operator is often referred to as *generalised projection*.

Computation of proximal points: examples

Fix for now $\gamma = 1$.

- (**Constant**) If $g(x) = c$, $c \in \mathbb{R}$ for every x (constant function). Then:

$$\text{prox}_g(x) = \arg \min_{y \in \mathbb{R}^n} c + \frac{1}{2} \|x - y\|^2 = x$$

Computation of proximal points: examples

Fix for now $\gamma = 1$.

- (Constant) If $g(x) = c$, $c \in \mathbb{R}$ for every x (constant function). Then:

$$\text{prox}_g(x) = \arg \min_{y \in \mathbb{R}^n} c + \frac{1}{2} \|x - y\|^2 = x$$

- (Affine) If $g(x) = a^T x + b$, for $a \in \mathbb{R}^n$ and $b \in \mathbb{R}$. Then:

$$\begin{aligned} \text{prox}_g(x) &= \arg \min_{y \in \mathbb{R}^n} a^T y + b + \frac{1}{2} \|y - x\|^2 \\ &= \arg \min_{y \in \mathbb{R}^n} a^T x + b - \frac{1}{2} \|a\|^2 + \frac{1}{2} \|y - (x - a)\|^2 = x - a \quad (\text{translation}) \end{aligned}$$

Computation of proximal points: examples

Fix for now $\gamma = 1$.

- (Constant) If $g(x) = c$, $c \in \mathbb{R}$ for every x (constant function). Then:

$$\text{prox}_g(x) = \arg \min_{y \in \mathbb{R}^n} c + \frac{1}{2} \|x - y\|^2 = x$$

- (Affine) If $g(x) = a^T x + b$, for $a \in \mathbb{R}^n$ and $b \in \mathbb{R}$. Then:

$$\begin{aligned} \text{prox}_g(x) &= \arg \min_{y \in \mathbb{R}^n} a^T y + b + \frac{1}{2} \|y - x\|^2 \\ &= \arg \min_{y \in \mathbb{R}^n} a^T x + b - \frac{1}{2} \|a\|^2 + \frac{1}{2} \|y - (x - a)\|^2 = x - a \quad (\text{translation}) \end{aligned}$$

- (Quadratic) If $g(x) = \frac{1}{2} x^T A x + b^T x + c$ with $A \in \mathbb{R}^{n \times n}$ and SPD, $b \in \mathbb{R}^n$ and $c \in \mathbb{R}$. Then:

$$\text{prox}_g(x) = \arg \min_{y \in \mathbb{R}^n} \frac{1}{2} y^T A y + b^T y + c + \frac{1}{2} \|y - x\|^2$$

Take the gradient and set it to 0:

$$A\hat{y} + b + \hat{y} - x = 0 \quad \Rightarrow \quad (A + \text{Id})\hat{y} = x - b \quad \Rightarrow \quad \hat{y} = (A + \text{Id})^{-1}(x - b)$$

Computation of proximal points: properties

Exercise: compute, for $\tau > 0$ $\text{prox}_{\tau g}(x)$ where $g(x) = |x|$. Plot the result as a function of x .

Computation of proximal points: properties

Exercise: compute, for $\tau > 0$ $\text{prox}_{\tau g}(x)$ where $g(x) = |x|$. Plot the result as a function of x .

Proposition (prox of separable functions)

Let $g : \mathbb{R}^n \rightarrow \mathbb{R} \cup \{+\infty\}$ be proper, convex, l.s.c. and **separable**, with $g(x) = \sum_{i=1}^n g_i(x_i)$ for proper, convex, l.s.c. 1D functions $g_i : \mathbb{R} \rightarrow \mathbb{R} \cup \{+\infty\}$. Then for $\gamma > 0$

$$\text{prox}_{\gamma g}(x) = (\text{prox}_{\gamma g_1}(x_1), \dots, \text{prox}_{\gamma g_n}(x_n)),$$

so the prox of a multi-dimensional function can be computed as the vector of prox's of their components g_i .

Exercise: For $\tau > 0$, give the expression of the proximal operator $\text{prox}_{\tau g}(x)$ with $g(x) = \|x\|_1$ for $x \in \mathbb{R}^n$.

Exercise: Behaviour of prox w.r.t. scaling/quadratic perturbations.

Projected gradient descent

For proper, differentiable, convex f and convex, closed $C \in \mathbb{R}^n$:

$$\arg \min_{x \in C} f(x) = \arg \min_{x \in \mathbb{R}^n} f(x) + \iota_C(x)$$

Algorithm (PGD)

Input: $x_0 \in \mathbb{R}^n$ (initial guess), $\tau \in (0, \frac{2}{L})$ (step-size)

Iterate for $k \geq 0$:

$$x_{k+\frac{1}{2}} = x_k - \tau \nabla f(x_k)$$

$$\begin{aligned} x_{k+1} &= P_C(x_{k+\frac{1}{2}}) = \arg \min_{y \in C} \frac{1}{2} \|y - x_{k+\frac{1}{2}}\|^2 \\ &= \arg \min_{y \in \mathbb{R}^n} \iota_C(y) + \frac{1}{2} \|y - x_{k+\frac{1}{2}}\|^2 = \text{prox}_{\iota_C}(x_{k+\frac{1}{2}}) \end{aligned}$$

till convergence.

$\Rightarrow$ the algorithm converges to a minimiser of f

- First: **gradient step**, next **projection step**
- Note: the projection on C is seen as an **inner minimisation problem** with explicit solution

Forward-backward splitting

a.k.a. proximal gradient descent

Why all this?

So far, we have discussed:

- **gradient descent** for minimising proper convex, differentiable functions f
- **implicit gradient descent** (proximal operators) for minimising proper convex (non-differentiable) functions g

Idea: combine the two ideas for solving the original, composite problem

$$\min_{x \in \mathbb{R}^n} \{F(x) := f(x) + g(x)\},$$

Why all this?

So far, we have discussed:

- **gradient descent** for minimising proper convex, differentiable functions f
- **implicit gradient descent** (proximal operators) for minimising proper convex (non-differentiable) functions g

Idea: combine the two ideas for solving the original, composite problem

$$\min_{x \in \mathbb{R}^n} \{F(x) := f(x) + g(x)\},$$

Forward-backward splitting (FB/FBS) algorithm

Input: $x_0 \in \mathbb{R}^n$, $\tau \in (0, \frac{2}{L})$

For $k \geq 0$, iterate:

$$x_{k+1} = \text{prox}_{\tau g}(x_k - \tau \nabla f(x_k)) \Leftrightarrow \begin{cases} x_{k+1/2} &= x_k - \tau \nabla f(x_k) \\ x_{k+1} &= \text{prox}_{\tau g}(x_{k+1/2}) \end{cases}$$

till convergence.

Such scheme alternates explicit (**forward**) and implicit (**backward**) gradient descent for minimising f and g alternatively. It's called **forward-backward** splitting algorithm or **proximal gradient** algorithm (Combettes, Wajs, '15).

Main ingredients:

- **Nice (smooth) part:** proper, convex, differentiable function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ with L -Lipschitz continuous gradient

$$(\forall x, y \in \mathbb{R}^n) \quad \|\nabla f(x) - \nabla f(y)\| \leq L\|x - y\|$$

- **Non-nice (non-smooth) part:** proper, l.s.c., convex (non-smooth) function $g : \mathbb{R}^n \rightarrow \mathbb{R} \cup \{+\infty\}$

Typically, we assume that g is *easily proximable*, i.e. the proximal operator of g can be computed in closed form (examples are $g(x) = \iota_C(x)$, $g(x) = \|x\|_1$, $g(x) = \|x\|_1 + \iota_{\geq 0}(x)$, $g(x) = \|x\|_1 + \frac{\lambda}{2}\|x\|_2^2 \dots$).

Forward-backward splitting (FB/FBS) algorithm¹

Input: $x_0 \in \mathbb{R}^n$, $\tau \in (0, \frac{2}{L})$

For $k \geq 0$, iterate:

$$x_{k+1} = \text{prox}_{\tau g}(x_k - \tau \nabla f(x_k))$$

till convergence.

- Step-size: same bound depending on L as for GD.
- If g is easily proximable: no inner minimisation. Otherwise: need to solve a nested minimisation problem up to some accuracy (**inexact** algorithms)
- Computational cost/complexity: evaluation of ∇f may be costly (matrix/vector products), number of iterations before convergence depends on τ .
 - * Too small τ : unnecessary too many iterations
 - * Too big τ : risk of moving to a point z for which $f(x) > f(x_k) \dots$

¹Combettes, Wajs, 2005

- If $g \equiv 0$: smooth-optimisation problem. The FB scheme is nothing but **GD**.
- If $g(x) = \iota_C(x)$ for closed and convex $C \rightarrow$ **projected gradient descent**.
- If $g(x) = \lambda\|x\|_1$ for $\lambda > 0$, the problem reads

$$\min_{x \in \mathbb{R}^n} f(x) + \lambda\|x\|_1, \quad (*)$$

then the algorithm to solve it is typically known as **ISTA**.

²Daubechies, Defrise, De Mol, 2004

- If $g \equiv 0$: smooth-optimisation problem. The FB scheme is nothing but **GD**.
- If $g(x) = \iota_C(x)$ for closed and convex $C \rightarrow$ **projected gradient descent**.
- If $g(x) = \lambda\|x\|_1$ for $\lambda > 0$, the problem reads

$$\min_{x \in \mathbb{R}^n} f(x) + \lambda\|x\|_1, \quad (*)$$

then the algorithm to solve it is typically known as **ISTA**.

Iterative Soft Thresholding Algorithm (ISTA)²

The iteration (*) takes the form:

$$x_{k+1} = \mathcal{T}_{\tau\lambda}(x_k - \tau\nabla f(x_k))$$

where for $z = x_k - \tau\nabla f(x_k)$, $\mathcal{T}_{\tau\lambda}(z)$ is the *soft-thresholding* (or *shrinkage*) operator (prox of the function $g(x) = \lambda\|x\|_1$ with parameter τ) defined by:

$$\mathcal{T}_{\tau\lambda}(z) = (\mathcal{T}_{\tau\lambda}(z_j))_{j=1,\dots,n} = \left([|z_j| - \lambda\tau]_+ \operatorname{sign}(z_j) \right)_{j=1,\dots,n}$$

²Daubechies, Defrise, De Mol, 2004

Convergence properties of proximal gradient descent

Convergence rate of FBS/PGD

Let $x^0 \in \mathbb{R}^n$, and $(x_k)_k$ the sequence of iterates generated by the FB algorithm. If $\tau \in (0, 2/L)$, there holds:

$$\underbrace{F(x_k)}_{(f+g)(x_k)} - F(x^*) \leq \underbrace{\frac{\|x_0 - x^*\|^2}{2\tau}}_{C(x^0, x^*, \tau) :=} \frac{1}{k} = O\left(\frac{1}{k}\right),$$

where x^* is a minimiser of F .

Convergence rate of FBS/PGD

Let $x^0 \in \mathbb{R}^n$, and $(x_k)_k$ the sequence of iterates generated by the FB algorithm. If $\tau \in (0, 2/L)$, there holds:

$$\underbrace{F(x_k)}_{(f+g)(x_k)} - F(x^*) \leq \underbrace{\frac{\|x_0 - x^*\|^2}{2\tau}}_{C(x^0, x^*, \tau) :=} \frac{1}{k} = O\left(\frac{1}{k}\right),$$

where x^* is a minimiser of F .

- The point x^* is unknown, so the constant C is.
- The difference between x_k and x^* is measured in function values: the values $F(x_k)$ asymptotically approaches the (unknown) value $F(x^*)$ with speed proportional to $1/k$
- The result can also be written $O(1/k)$ where the constant is unknown, as $\|x_0 - x^*\|^2$ is.

Accelerated forward-backward splitting (FISTA)

Accelerated forward-backward algorithm (FISTA)

Analogous idea as for GD: combine proximal updates with inertia.

Accelerated forward-backward splitting (FISTA)

Input: $x_0 = x^{-1} = y_0 \in \mathbb{R}^n$ (initialisations), $\tau \in (0, \frac{1}{L}]$ (step-size), $t_0 = 1$

Iterate for $k \geq 0$:

$$y_k = x_k + \frac{t_k - 1}{t_{k+1}}(x_k - x_{k-1})$$

$$x_{k+1} = \text{prox}_{\tau g}(y_k - \tau \nabla f(y_k))$$

$$t_{k+1} = \frac{1 + \sqrt{1 + 4t_k^2}}{2}$$

till convergence.


- Proximal operator is evaluated in GD-type applied on the extrapolated sequence (y_k)
- Same computational costs as FB: one gradient computation and one prox computation at each iteration + explicit update. Modifications are cheap!

Convergence rates for FISTA

Convergence rate of FISTA

For a convex, smooth function f with L -Lipschitz gradient and convex, proper and l.s.c. function g , the convergence rate for the function values of FISTA for the composite function $F = f + g$ with $\tau \in (0, 1/L]$ is:

$$F(x_k) - F(x^*) \leq \frac{1}{(k+1)^2} \frac{2\|x_0 - x^*\|^2}{\tau}$$



Questions?

calatroni@i3s.unice.fr