

Introduction to Reinforcement Learning

4/10

Jean Martinet

MSc DSAI

2024 – 2025

- Introduction
 - Course 1 : Introduction to Reinforcement Learning (RL)
- Part I on tabular methods
 - Course 2 : Markov Decision Processes
 - Course 3 : Dynamic programming in RL
 - **Course 4 : Temporal difference 1/2 (Q-learning)**
 - Course 5 : Temporal difference 2/2 (SARSA)
- Part II on approximate methods
 - Course 6 : Value function approximation
 - Course 7 : Eligibility traces
 - Course 8 : Policy gradient 1/2 (REINFORCE)
 - Course 9 : Policy gradient 2/2 (actor-critic methods)
 - Course 10 : Projects presentation session

Reminder : think of a project topic

- **Choose from :**

- Public presentation of articles/advanced topics/applications
 - Conference paper or book chapter
 - Advanced theme (e.g. actor-critic, eligibility trace, etc.)
 - Application domain (e.g. temperature control, revenue management, etc.)
- Deepening or exploration project
 - Subject to be chosen/defined and validated

- **Choice to be validated... today**

- **Expected result :**

- Short 2-page max PDF report
- Code (ipynb / py / git)
- Short 10-min presentation during last / before last session

About the project

- Double objective
 - Dig deeper in a specific subject (discussed or not during the lectures)
 - Share your insights with other students (in a teacher mode)
- A bit hard to choose early, before having reviewed all topics
- If you can define what is the environment, the reward, the agent, and the actions, it is a good start
- Stay small, at least for a first version, then make it more complex if you have time
- An experimental contribution is needed
 - E.g. compare two algorithms
 - E.g. start from an existing approach, and monitor changes when parameters vary
- The project needs be ORIGINAL
 - You need an original contribution of your own
 - Make sure your project is different from what can be found online
- IMPORTANT : if you decide to use an existing work, it is MANDATORY to cite the source, and you need to state what your contribution is

Today's menu

- Monte Carlo methods
- Dynamic programming vs Temporal difference
- On-policy vs off-policy learning
- Exploration vs exploitation
- Optimal policy approximation with Q-learning

- ✓ MC does not assume a complete knowledge of the environment
 - MC methods require only *experience*, contrary to DP
 - A model is required
 - Obtain optimal behaviour without any prior information about environment dynamics
 - It is sometimes easier to generate transitions than to obtain complete probability distributions in explicit forms
- The method is based on averaging sample returns
 - Strong assumption : episodic tasks only – to insure that well defined returns are available
 - Only after episode completion, value estimates and policies are changed – not step-by-step updates

MC prediction

- Prediction means : "evaluate policies"
- Simple average of returns observed after visits to each state
- Converges to the true expected values when the number of visits to s goes to infinity
- Two-versions : *first-visit* and *every-visit*

First-visit MC prediction, for estimating $V \approx v_\pi$

Input: a policy π to be evaluated

Initialize:

$V(s) \in \mathbb{R}$, arbitrarily, for all $s \in \mathcal{S}$

$Returns(s) \leftarrow$ an empty list, for all $s \in \mathcal{S}$

✓ Loop forever (for each episode):

Generate an episode following π : $S_0, A_0, R_1, S_1, A_1, R_2, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

✓ Loop for each step of episode, $t = T-1, T-2, \dots, 0$:

$G \leftarrow \gamma G + R_{t+1}$

Unless S_t appears in $S_0, S_1, \dots, S_{t-1}$:


Append G to $Returns(S_t)$

$V(S_t) \leftarrow \text{average}(Returns(S_t))$

Evaluation of action values

- With a model, state values alone are sufficient to determine a policy
 - Just select the action that leads to the best (s_{t+1}, r_{t+1}) as in DP
- When a model is not available, we need to explicitly evaluate the value of actions
 - ✓ Estimate $Q^\pi(s, a)$ similarly, except that we now deal with (s, a) visits
- However, many (s, a) pairs may never be visited if π is deterministic
 - Option to *force* the episode to start with all possible (s, a)
 - This is the *exploring starts* assumption – in simulation not real world

- Control means : "approximate optimal policies"
- Generalised Policy iteration
 - Greedy policy with the same pattern as DP with $\pi(s) \doteq \arg \max_a Q^\pi(s, a)$



$$\pi_0 \xrightarrow{\text{E}} q_{\pi_0} \xrightarrow{\text{I}} \pi_1 \xrightarrow{\text{E}} q_{\pi_1} \xrightarrow{\text{I}} \pi_2 \xrightarrow{\text{E}} \dots \xrightarrow{\text{I}} \pi_* \xrightarrow{\text{E}} q_*$$

- Simplest TD update

$$V(s_t) \leftarrow V(s_t) + \alpha[r_{t+1} + \gamma V(s_{t+1}) - V(s_t)]$$

- The amount in brackets is an error measuring the difference between
 - the estimated value of S_t and
 - the better estimate $r_{t+1} + \gamma V(s_{t+1}) - V(s_t)$
- This quantity is called the *TD error*

$$\delta_t \doteq r_{t+1} + \gamma V(s_{t+1}) - V(s_t)$$



Tabular TD(0) for estimating v_π

Input: the policy π to be evaluated

Algorithm parameter: step size $\alpha \in (0, 1]$

Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

$A \leftarrow$ action given by π for S

 Take action A , observe R, S'

$V(S) \leftarrow V(S) + \alpha [R + \gamma V(S') - V(S)]$

$S \leftarrow S'$


 until S is terminal



Dynamic programming vs Time difference

- In DP, we assumed a complete knowledge of the environment (MDP)
 - We computed value functions
 - Now with TD, we learn them
- We require only experience
 - Actual : no prior knowledge of environment is required
 - Simulated : a model is required to generate sample transitions
 - (it is easier to get sample transitions than the complete distribution)
- Everything else remains similar

On-policy vs off-policy

- 
- Our agent learns action values relying on our current policy estimate
 - But it needs to behave non-optimally to explore all actions
 - (and find optimal actions)
 - We can use – and distinguish – two policies
 - The policy that we learn and becomes optimal (target π)
 - The exploratory policy that generates behaviour (behaviour π)
 - This is *off-policy* learning
 - On the contrary, *on-policy* methods use a single policy

Greedy actions vs exploratory actions



- Simplest action selection rule : select one of the actions with the highest estimated value

$$A_t \doteq \arg \max_a Q^\pi(s, a)$$

- Greedy action selection exploit current knowledge to maximise immediate reward
- Alternative :
 - Behave greedily most of the time
 - Sometimes, with probability ϵ , randomly pick another action (exploratory move)
- This is the ϵ -greedy method

- Trade-off between exploration and exploitation using ϵ value
- ✓ • ϵ starts at 1 (only exploration) then decreases (more exploitation)
- Choose a random number r between 0 and 1 (uniform distribution) :
 - if $r < \epsilon$: choose a random action (exploration)
 - if $r \geq \epsilon$: choose the best action = maximising Q value (exploitation)

Q-learning algorithm

- Here again, π evaluation (prediction) and π improvement towards π^* (control)
- Q-learning is an off-policy TD algorithm

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha[r_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t)]$$


- Iteratively updates action-value using δ_t
- α is the learning rate
- Build a Q-table of size $|S| \times |A|$, initial values are 0s

- Iterate episodes (a few hundreds)

- Iterate steps in each episode

- Select the best action a in state s , use the reward to update Q
- Episode terminates when agent reaches a terminal state (or max iteration)

Today's lab (graded, duration 1h30)



- 1. Consider a 1D grid with :
 - one goal location (positive reward, e.g. +1)
 - one trap location (negative reward, e.g. -1)
 - a fixed move cost (e.g. -0.01)
 - deterministic actions (probability to go left when trying left is 1)
- Implement Q-learning from the equation, with ϵ -greedy
- Run your algorithm to determine the best policy, check it
- 2. Extend to the 2D grid of the classical toy example (lecture 1)
- 3. Optional : Visualise results with plots for several values of γ , α , and ϵ (brings bonus!)
- (Note : use Python and numpy only – gym next week)
- Submit your solution after 90 min